


Fig 11 — ITL Tail Latency CDF under BurstGPT Traffic
(vLLM, Llama-2-7B, tensor_parallel=4, gpu_util=0.65)



(synthetic — run phase0/run_itl_cdf_eval.slurm for real data)